

Apache Spark Structured API

- The Structured APIs are a tool for manipulating all sorts of data.
- Unstructured log files to semi-structured CSV files.
- Highly structured Parquet files.
- These APIs refer to the following core types of distributed collection:
- SQL tables and views, DataFrames and Datasets

Spark SQL

- You can run SQL queries against views or tables organized into databases
- You also can use system functions or define user functions and analyze query plans in order to optimize their workloads.
- This integrates directly into the DataFrame and Dataset API
- You can choose to express some of your data manipulations in SQL and others in DataFrames

```
In [1]: # Import dependencies
from pyspark.sql import SparkSession
from pyspark.sql import *
from pyspark.sql.functions import *
from pyspark.sql.types import *
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
spark = SparkSession(sc)
```

Creating Tables

- CREATE and DROP tables
- Create table from .json file
- Create table from .csv file

```
In [2]: # Drop table flights
spark.sql("""
DROP TABLE IF EXISTS flights
""")
```

Out[2]: DataFrame[]

```
In [3]: # Create table 'flights' with columns
# DEST_COUNTRY_NAME STRING, ORIGIN_COUNTRY_NAME STRING, count LONG
# From .json file: path Resources/20015-summary.json
spark.sql("""
CREATE TABLE flights (
DEST_COUNTRY_NAME STRING, ORIGIN_COUNTRY_NAME STRING, count LONG)
USING JSON OPTIONS (path '/Resources/2015-summary.json')
""")
```

Out[3]: DataFrame[]

```
In [4]: # Drop flights_csv
spark.sql("""
DROP TABLE IF EXISTS flights_csv
""")
```

Out[4]: DataFrame[]

```
In [5]: # Create table 'flights_csv' with columns
# DEST_COUNTRY_NAME STRING, ORIGIN_COUNTRY_NAME STRING, count LONG
# From .csv file 20015-summary.csv
spark.sql("""
CREATE TABLE flights_csv (
DEST_COUNTRY_NAME STRING,
ORIGIN_COUNTRY_NAME STRING COMMENT "remember, the US will be most prevalent",
count LONG)
USING csv OPTIONS (header true, path '/Resources/2015-summary.csv')
""")
```

Out[5]: DataFrame[]

Views

- Creating Views
- Creating Temporary views that are available only during the current session
- Overwrite and replace view if one already exists from previously

```
In [6]: # We create a view(just_usa_view) in which the dest_country_name is United States in ord
spark.sql("""
CREATE VIEW just_usa_view AS
SELECT * FROM flights WHERE dest_country_name = 'United States'
""")
```

Out[6]: DataFrame[]

```
In [7]: # Create a temporary view(just_usa_view) in which the destination is United States in o
spark.sql("""
CREATE TEMP VIEW just_usa_view_temp AS
SELECT * FROM flights WHERE dest_country_name = 'United States'
""")
```

Out[7]: DataFrame[]

```
In [8]: # Overwrite and replace view(just_usa_view) if one already exists from previous
spark.sql("""
CREATE OR REPLACE TEMP VIEW just_usa_view_temp AS
SELECT * FROM flights WHERE dest_country_name = 'United States'
""")
```

Out[8]: DataFrame[]

DataFrame Transformations

- Adding/removing rows
- Transform row into column(or vice versa)
- Changing the order of rows based on the values in columns

Creating DataFrames

```
In [9]: # Create a new DataFrame and register as a temporary view to query it in SQL(SQL trasfo
# For Querying in SQL, name the SQL dataset table (2015-summary-json)
df = spark.read.format("json").load("Resources/2015-summary.json")
df.createOrReplaceTempView("dfTable")
```

```
In [10]: # Return the schema(StructType) of the DataFrame: Schemas define the name as well as th
# Return the the first five column records
# Return statistics for numeric columns
# Return the logical and physical plans. DataFrame lineage(how Spark executes query)
# Return the datatypes
df.printSchema()
df.show(5)
df.describe().show()
df.explain()
df.dtypes
```

root

```
-- DEST_COUNTRY_NAME: string (nullable = true)
-- ORIGIN_COUNTRY_NAME: string (nullable = true)
-- count: long (nullable = true)
```

```
+-----+-----+-----+
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|
+-----+-----+-----+
|    United States|          Romania|    15|
|    United States|          Croatia|     1|
|    United States|          Ireland|   344|
|         Egypt|    United States|    15|
|    United States|          India|    62|
+-----+-----+-----+
```

only showing top 5 rows

```
+-----+-----+-----+-----+
|summary|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|
+-----+-----+-----+-----+
|count|256|256|256|
|mean|null|null|1770.765625|
|stddev|null|null|23126.516918551915|
|min|Algeria|Angola|1|
|max|Zambia|Vietnam|370002|
+-----+-----+-----+-----+
```

== Physical Plan ==

```
*(1) FileScan json [DEST_COUNTRY_NAME#9,ORIGIN_COUNTRY_NAME#10,count#11L] Batched: fals
e, Format: JSON, Location: InMemoryFileIndex[file:/C:/Users/tenle/Documents/Web/Spark AP
I Structured Operations/Resources/20..., PartitionFilters: [], PushedFilters: [], ReadSc
hema: struct<DEST_COUNTRY_NAME:string,ORIGIN_COUNTRY_NAME:string,count:bigint>
```

```
Out[10]: [('DEST_COUNTRY_NAME', 'string'),
('ORIGIN_COUNTRY_NAME', 'string'),
('count', 'bigint')]
```

```
In [11]: # From the df table return the DES_COUNTRY_Name column(return 2 rows)
df.selectExpr("DEST_COUNTRY_NAME").show(2)
```

```
+-----+
|DEST_COUNTRY_NAME|
+-----+
|    United States|
|    United States|
+-----+
only showing top 2 rows
```

Adding, Renaming and Dropping Columns

- using as [column_name]
- using withColumn method
- using the withColumnRenamed method

```
In [12]: # Add a new column withinCountry to our DataFrame that specifies whether the destination
df.selectExpr(
    "*", # all original columns
    "(DEST_COUNTRY_NAME = ORIGIN_COUNTRY_NAME) as withinCountry")\
    .show(2)
```

```
+-----+-----+-----+-----+
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|withinCountry|
+-----+-----+-----+-----+
|    United States|          Romania|    15|         false|
|    United States|          Croatia|     1|         false|
+-----+-----+-----+-----+
only showing top 2 rows
```

```
In [13]: # Add a column name [numberOne] with values of 1, us withColumn method
df.withColumn("numberOne", lit(1)).show(2)
```

```
+-----+-----+-----+-----+
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|numberOne|
+-----+-----+-----+-----+
|    United States|          Romania|    15|         1|
|    United States|          Croatia|     1|         1|
+-----+-----+-----+-----+
only showing top 2 rows
```

```
In [14]: # set a Boolean flag for when the origin country is the same as the destination country
df.withColumn("withinCountry", expr("ORIGIN_COUNTRY_NAME == DEST_COUNTRY_NAME"))\
    .show(2)
```

```
+-----+-----+-----+-----+
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|withinCountry|
+-----+-----+-----+-----+
|    United States|          Romania|    15|         false|
|    United States|          Croatia|     1|         false|
+-----+-----+-----+-----+
only showing top 2 rows
```

```
In [15]: # Rename the DEST_COUNTRY_NAME column to desc
df.withColumnRenamed("DEST_COUNTRY_NAME", "dest").columns
```

```
Out[15]: ['dest', 'ORIGIN_COUNTRY_NAME', 'count']
```

```
In [16]: # Remove the ORIGIN_COUNTRY_NAME columns
df.drop("ORIGIN_COUNTRY_NAME").columns
```

```
Out[16]: ['DEST_COUNTRY_NAME', 'count']
```

Changing a Column's Type (cast)

```
In [17]: # Python Let's convert our count column from an integer to a type Long:
df.withColumn("count2", col("count").cast("long"))
```

```
Out[17]: DataFrame[DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string, count: bigint, count2:
bigint]
```

```
In [18]: # SQL Let's convert our count column from an integer to a type Long:
spark.sql("""
SELECT *, cast(count as long) AS count2 FROM dfTable
""").show(3)
```

DEST_COUNTRY_NAME	ORIGIN_COUNTRY_NAME	count	count2
United States	Romania	15	15
United States	Croatia	1	1
United States	Ireland	344	344

only showing top 3 rows

Filtering & Sorting Rows

```
In [19]: # Python - filter the count column returning values < 2
# Origin_country_name is not equal to Croatia, show two rows
df.where(col("count") < 2).where(col("ORIGIN_COUNTRY_NAME") != "Croatia")\
.show(2)
```

DEST_COUNTRY_NAME	ORIGIN_COUNTRY_NAME	count
United States	Singapore	1
Moldova	United States	1

only showing top 2 rows

```
In [20]: # SQL - filter the count column returning values < 2
# Origin_country_name is not equal to Croatia, show two rows
spark.sql("""
SELECT *
FROM dfTable
WHERE count < 2
AND ORIGIN_COUNTRY_NAME != "Croatia"
LIMIT 2
""").show()
```

DEST_COUNTRY_NAME	ORIGIN_COUNTRY_NAME	count
United States	Singapore	1
Moldova	United States	1

```
+-----+-----+-----+
```

```
In [21]: # Return the number of Unique rows from ORIGIN_COUNTRY_NAME
df.select("ORIGIN_COUNTRY_NAME", "DEST_COUNTRY_NAME").distinct().count()
```

```
Out[21]: 256
```

```
In [22]: spark.sql("""
SELECT COUNT(DISTINCT(ORIGIN_COUNTRY_NAME, DEST_COUNTRY_NAME))
FROM dfTable
""").show()
```

```
+-----+-----+-----+
|count(DISTINCT named_struct(ORIGIN_COUNTRY_NAME, ORIGIN_COUNTRY_NAME, DEST_COUNTRY_NAME)|
+-----+-----+-----+
|
256|
+-----+-----+-----+
```

```
In [23]: # Python - Return the 'count' column in 'desc order' and the DEST_COUNTRY_NAME column i
df.orderBy(col("count").desc(), col("DEST_COUNTRY_NAME").asc()).show(2)
```

```
+-----+-----+-----+
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME| count|
+-----+-----+-----+
|    United States|    United States|370002|
|    United States|          Canada| 8483|
+-----+-----+-----+
only showing top 2 rows
```